



Java Technologies

Messaging

The Context

- In a distributed system components interact with each other over a network, in order to achieve a common goal:
 - *Message Passing vs. Shared Memory*
- Message passing can be performed:
 - *Synchronous vs. Asynchronous*
- The components may interact:
 - directly with each other → *tight-coupled*
 - through an *intermediary* → *loose-coupled*

What is Messaging?

- Method of peer-to-peer communication between software components or applications, using an *intermediate agent*.
- **Loosely coupled, Asynchronous, Reliable**
 - differs from tightly coupled technologies, such as Remote Method Invocation (RMI), which require an application to know a remote application's methods.
 - differs from electronic mail (email), as it is used for communication between software applications or software components.

Use Cases

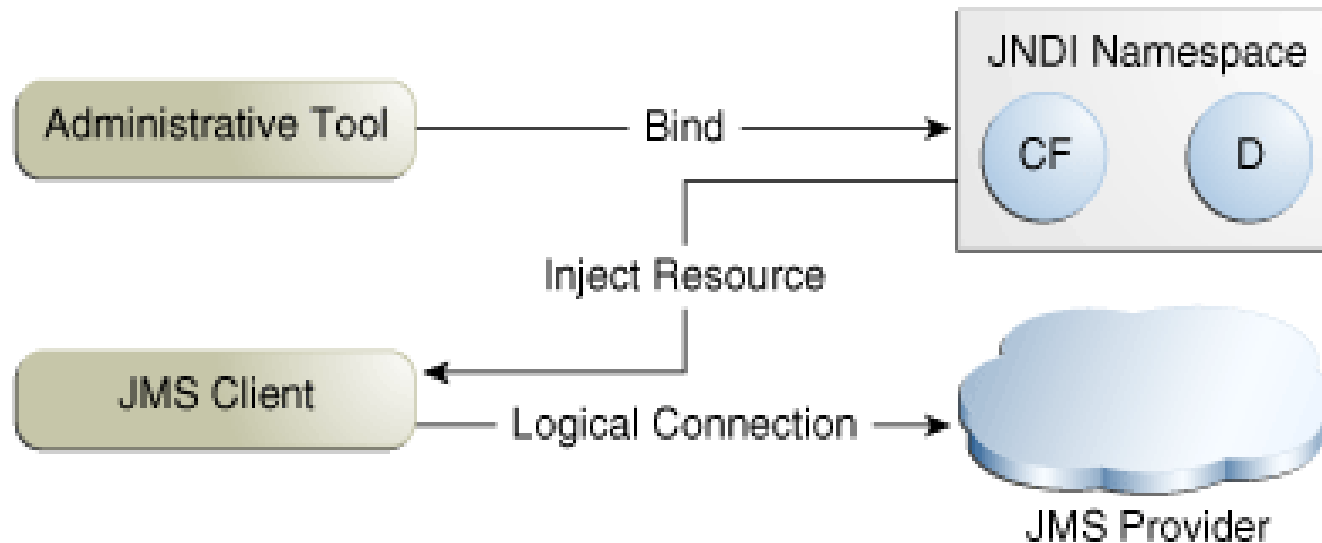
- **Loose coupling:** The provider wants the components not to depend on information about other components' interfaces, so components can be easily replaced
- **Fault-tolerance:** The provider wants the application to run whether or not all components are up and running simultaneously.
- **Asynchronous (delayed) communication:** The application business model allows a component to send information to another and to continue to operate without receiving an immediate response
- **Event-based communication**
- **Scalability:** Handling large number of operations

What is JMS?

Java Message Oriented Middleware (MOM) API


- The JMS API defines **a common set of interfaces** and associated semantics that allow programs written in the Java to communicate with other messaging implementations, regardless of the protocol.
- Allows applications to **create, send, receive, and read messages**.
- Offers **portability** of JMS applications across JMS providers.
- **Implementations:** OpenMQ (Oracle), ActiveMQ (Apache), RabbitMQ (Pivotal), JBoss Messaging, etc.
- Application servers usually have a JMS implementation included, by default.

JMS Architecture




- A **JMS provider** is a messaging system that implements the JMS interfaces and provides administrative and control features.
- **JMS clients** are the programs or components, that produce and consume messages. Any Java EE or SE application component can act as a JMS client.
- **Messages** are the objects that communicate information between JMS clients.
- **Administered objects** are JMS objects configured for the use of clients:
 - ✓ connection factories
 - ✓ destinations

Point-to-Point (PTP) Model



- PTP model is built on the concept of:
 - **message queues, senders, and receivers.**
- Each message is **addressed to a specific queue**, and receiving clients extract messages from the queues established to hold their messages.
- Queues retain all messages sent to them until the messages are consumed or expire.
- **Each message has only one consumer.**
- The receiver can fetch the message whether or not it was running when the client sent the message.

Publish/Subscribe Model

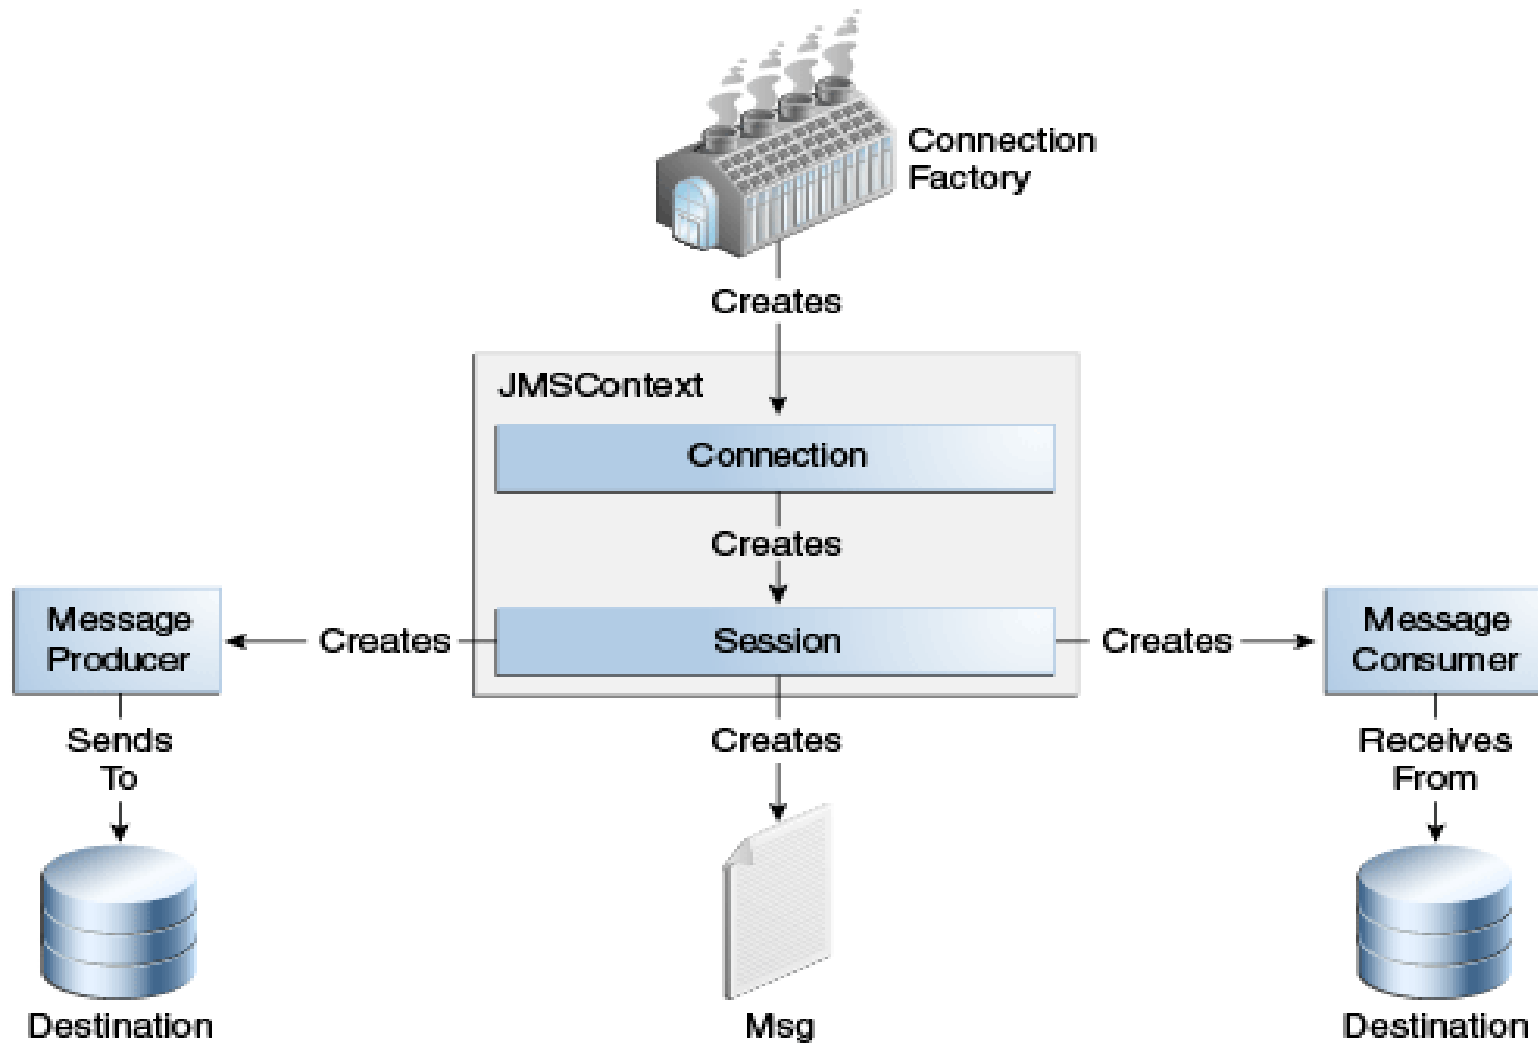


- Clients address messages to **a topic**.
- Publishers and subscribers can **dynamically publish or subscribe to the topic**.
- The system takes care of distributing the messages arriving from a topic's multiple publishers to its multiple subscribers. Topics retain messages only as long as it takes to distribute them to subscribers.
- **A topic can have many consumers**, but a subscription has **only one subscriber**.
- A client that subscribes to a topic can consume only messages sent after the client has created a subscription, and the consumer must continue to be active in order for it to consume messages (except for durable subscriptions).

Message Consumption

- **Synchronously / Blocking:** A consumer explicitly fetches the message from the destination by calling the *receive* method. The receive method can block until a message arrives or can time out if a message does not arrive within a specified time limit.
- **Asynchronously / Non-blocking:** An application client or a Java SE client can register a *message listener* with a consumer. calling the listener's *onMessage* method, which acts on the contents of the message.

The Programming Model



JMS Administered Objects

- A **connection factory** is the object a client uses to create a connection to a provider.

```
@Resource(lookup = "MyConnectionFactory")  
private ConnectionFactory connectionFactory;
```

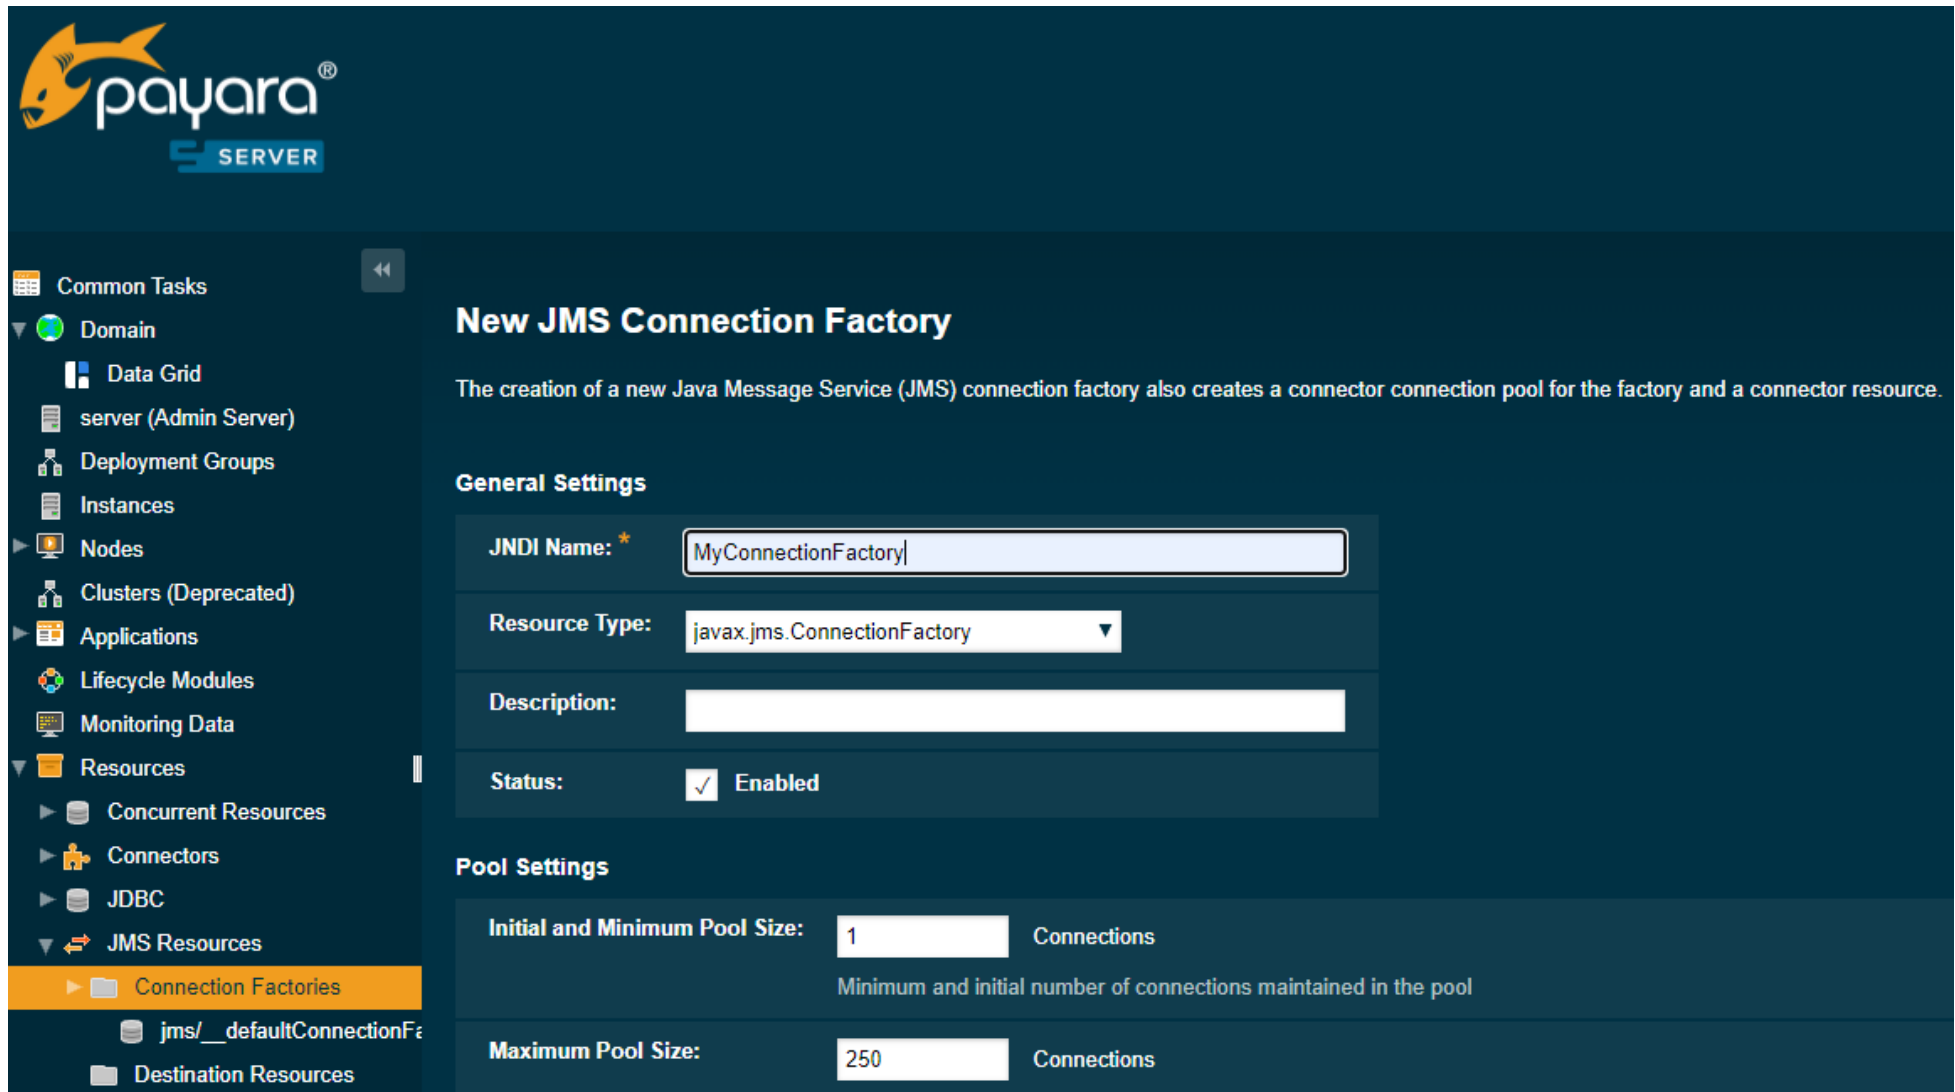
- A **destination** is the object a client uses to specify the target of messages it produces and the source of messages it consumes.

```
@Resource(lookup = "MyQueue")  
private Queue queue;
```

```
@Resource(lookup = "MyTopic")  
private Topic topic;
```

Example (Payara)

Payara already comes with OpenMQ



The screenshot displays the Payara Server administration console. The left sidebar contains a navigation menu with the following items: Common Tasks, Domain, Data Grid, server (Admin Server), Deployment Groups, Instances, Nodes, Clusters (Deprecated), Applications, Lifecycle Modules, Monitoring Data, Resources, Concurrent Resources, Connectors, JDBC, JMS Resources, Connection Factories (highlighted), jms/_defaultConnectionFactory, and Destination Resources. The main content area is titled 'New JMS Connection Factory' and includes a descriptive paragraph: 'The creation of a new Java Message Service (JMS) connection factory also creates a connector connection pool for the factory and a connector resource.' Below this, the 'General Settings' section contains the following fields: 'JNDI Name' (required, with a star icon) set to 'MyConnectionFactory', 'Resource Type' set to 'javax.jms.ConnectionFactory', 'Description' (empty), and 'Status' set to 'Enabled' with a checked checkbox. The 'Pool Settings' section contains two rows: 'Initial and Minimum Pool Size' set to '1' with a note 'Minimum and initial number of connections maintained in the pool', and 'Maximum Pool Size' set to '250'.

payara
SERVER

Common Tasks

Domain

Data Grid

server (Admin Server)

Deployment Groups

Instances

Nodes

Clusters (Deprecated)

Applications

Lifecycle Modules

Monitoring Data

Resources

Concurrent Resources

Connectors

JDBC

JMS Resources

Connection Factories

jms/_defaultConnectionFactory

Destination Resources

New JMS Connection Factory

The creation of a new Java Message Service (JMS) connection factory also creates a connector connection pool for the factory and a connector resource.

General Settings

JNDI Name: * MyConnectionFactory

Resource Type: javax.jms.ConnectionFactory

Description:

Status: ☒ Enabled

Pool Settings

Initial and Minimum Pool Size: 1 Connections
Minimum and initial number of connections maintained in the pool

Maximum Pool Size: 250 Connections

JMS Dependency

- Jakarta EE 8 Web Profile contains JMS
- Java EE 7 Web Profile does not contain JMS

```
<dependency>  
    <groupId>javax.jms</groupId>  
    <artifactId>javax.jms-api</artifactId>  
    <version>2.0.1</version>  
</dependency>
```

Connections, Sessions

- A **connection** encapsulates a virtual connection with a JMS provider. For example, a connection could represent an open TCP/IP socket between a client and a provider service daemon.
- You use a connection to create one or more sessions
- A **session** is a single-threaded context for producing and consuming messages.
- A session provides a **transactional context** with which to group a set of sends and receives into an atomic unit of work.

JMS Context

- A **JMSContext object** combines a *connection* and a *session* in a single object.
- Responsible with creation of JMS objects:
 - Message producers
 - Message consumers
 - Messages
 - Queue browsers

```
JMSContext context = connectionFactory.createContext();
```

Message Producers / Consumers

- A **message producer** is an object used for *sending messages to a destination*.
- A **message consumer** is an object used for *receiving messages sent to a destination*.

```
try(JMSContext context = connectionFactory.createContext();) {  
    JMSProducer producer = context.createProducer();  
    Message message = create the message ;  
    producer.send(dest, message); //dest is a queue or a topic  
    ...  
    JMSConsumer consumer = context.createConsumer(dest);  
    Message message = consumer.receive() ;  
    // or  
    consumer.setMessageListener( new MssageListener() {  
        public void onMessage(Message message) {  
            // Cod specific de tratare a mesajului  
        }  
    });  
}
```


Messages

A JMS message can have three parts:

- **Header:** values used by both clients and providers to identify and route messages:
 - JMSSessionId, JMSDestination, JMSType, JMSPriority, ...
- **Properties:** used to provide compatibility with other messaging systems, or to create *message selectors*.
- **Body:** [Text, Map, Bytes, Stream, Object] Message

```
TextMessage message = context.createTextMessage();  
message.setText("Hello World!");  
String body = message.getText();
```

Message Selectors

- A **selector** filters the messages it receives.
- It allows a message consumer to specify the messages it is interested in. Message selectors assign the work of filtering messages to the JMS provider rather than to the application.
- A message selector is a *String* that contains an expression. The syntax of the expression is based on a subset of the SQL92 conditional expression syntax. Examples:
 - `"NewsType = 'Sports' OR NewsType = 'Opinion'"`
 - `"JMSType = 'car' AND color = 'blue' AND weight > 2500"`
- A selector can be specified when you create a consumer.

The message consumer then receives only messages whose headers and properties match the selector. A message selector cannot select messages on the basis of the content of the message body.

Queue Browsers

- *Messages sent to a queue remain in the queue until the message consumer for that queue consumes them.*
- The JMS API provides a QueueBrowser object that allows you to browse the messages in the queue and display the header values for each message.

```
QueueBrowser browser = context.createBrowser(queue) ;
```

Acknowledgement & Transactions

- **Acknowledgement:** until a JMS message has been acknowledged, it is not considered to be successfully consumed.
 - AUTO_ACKNOWLEDGE
 - CLIENT_ACKNOWLEDGE → `message.acknowledge()`
- A **transaction** groups a series of operations into an *atomic unit of work*.

```
JMSContext context =  
    connectionFactory.createContext(JMSContext.SESSION_TRANSACTED);  
...  
context.commit(); // or context.rollback();
```

You can send multiple messages in a transaction, and *the messages will not be added to the queue or topic until the transaction is committed*. If you receive multiple messages in a transaction, *they will not be acknowledged until the transaction is committed*.

Sending Messages Asynchronously


- The send method *blocks* until the JMS provider confirms that the message was sent successfully.
- Sending a message asynchronously involves supplying a callback object → *CompletionListener*

```
CompletionListener listener = new SendListener();  
context.createProducer().setAsync(listener).  
    send(dest, message);
```

```
public class SendListener implements CompletionListener {  
    public void onCompletion(Message message) {  
        System.out.println("Send has completed.");  
    }  
    public void onException(Message message, Exception e) {  
        System.out.println("Send failed: " + e.toString());  
        System.out.println("Unsent message is: \n" + message);  
    }  
}
```

Message-Driven Beans

- Process messages **asynchronously**.
- Similar to an event listener, implements *javax.jms.MessageListener* interface.



Creating a MDB

@MessageDriven(mappedName="jms/Queue")

```
public class SimpleMessageBean implements MessageListener {  
    public void onMessage(Message inMessage) {  
        TextMessage msg = null;  
        try {  
            if (inMessage instanceof TextMessage) {  
                msg = (TextMessage) inMessage;  
            } else {  
                ...  
            }  
        } catch (JMSEException e) {  
            e.printStackTrace();  
        }  
    }  
}
```

Sending a Message

```
@Resource(mappedName="jms/ConnectionFactory")  
private ConnectionFactory connectionFactory;
```

```
@Resource(mappedName="jms/Queue") the destination  
private static Queue queue;
```

```
JMSContext context = connectionFactory.createContext();
```

```
producer = context.createProducer(queue);
```

```
Message message = context.createTextMessage();  
message.setText("Hello");
```

```
producer.send(message);
```


Advanced JMS Features

- Specifying Message Persistence
 - `DeliveryMode.PERSISTENT`, `NON-PERSISTENT`
- Setting Message Priority Levels
 - from 0 (lowest) to 9 (highest), the default level is 4
- Allowing Messages to Expire
 - `setTimeToLive`
- Specifying a Delivery Delay
 - `setDeliveryDelay`
- Creating Temporary Destinations

JMS Disadvantages

- JMS is a **Java API** for messaging. In applications using microservices this may be a problem.
- JMS is not a messaging protocol (and does not enforce one). If you want to learn about protocols take a look at:
 - **STOMP**: Simple (or Streaming) Text Orientated Messaging Protocol.
 - **AMQP**: Advanced Message Queuing Protocol
- Scalability can become an issues in complex scenarios. JMS may not be well suited to large scale message processing applications, where there is a requirement for very high throughput.

Event Streaming

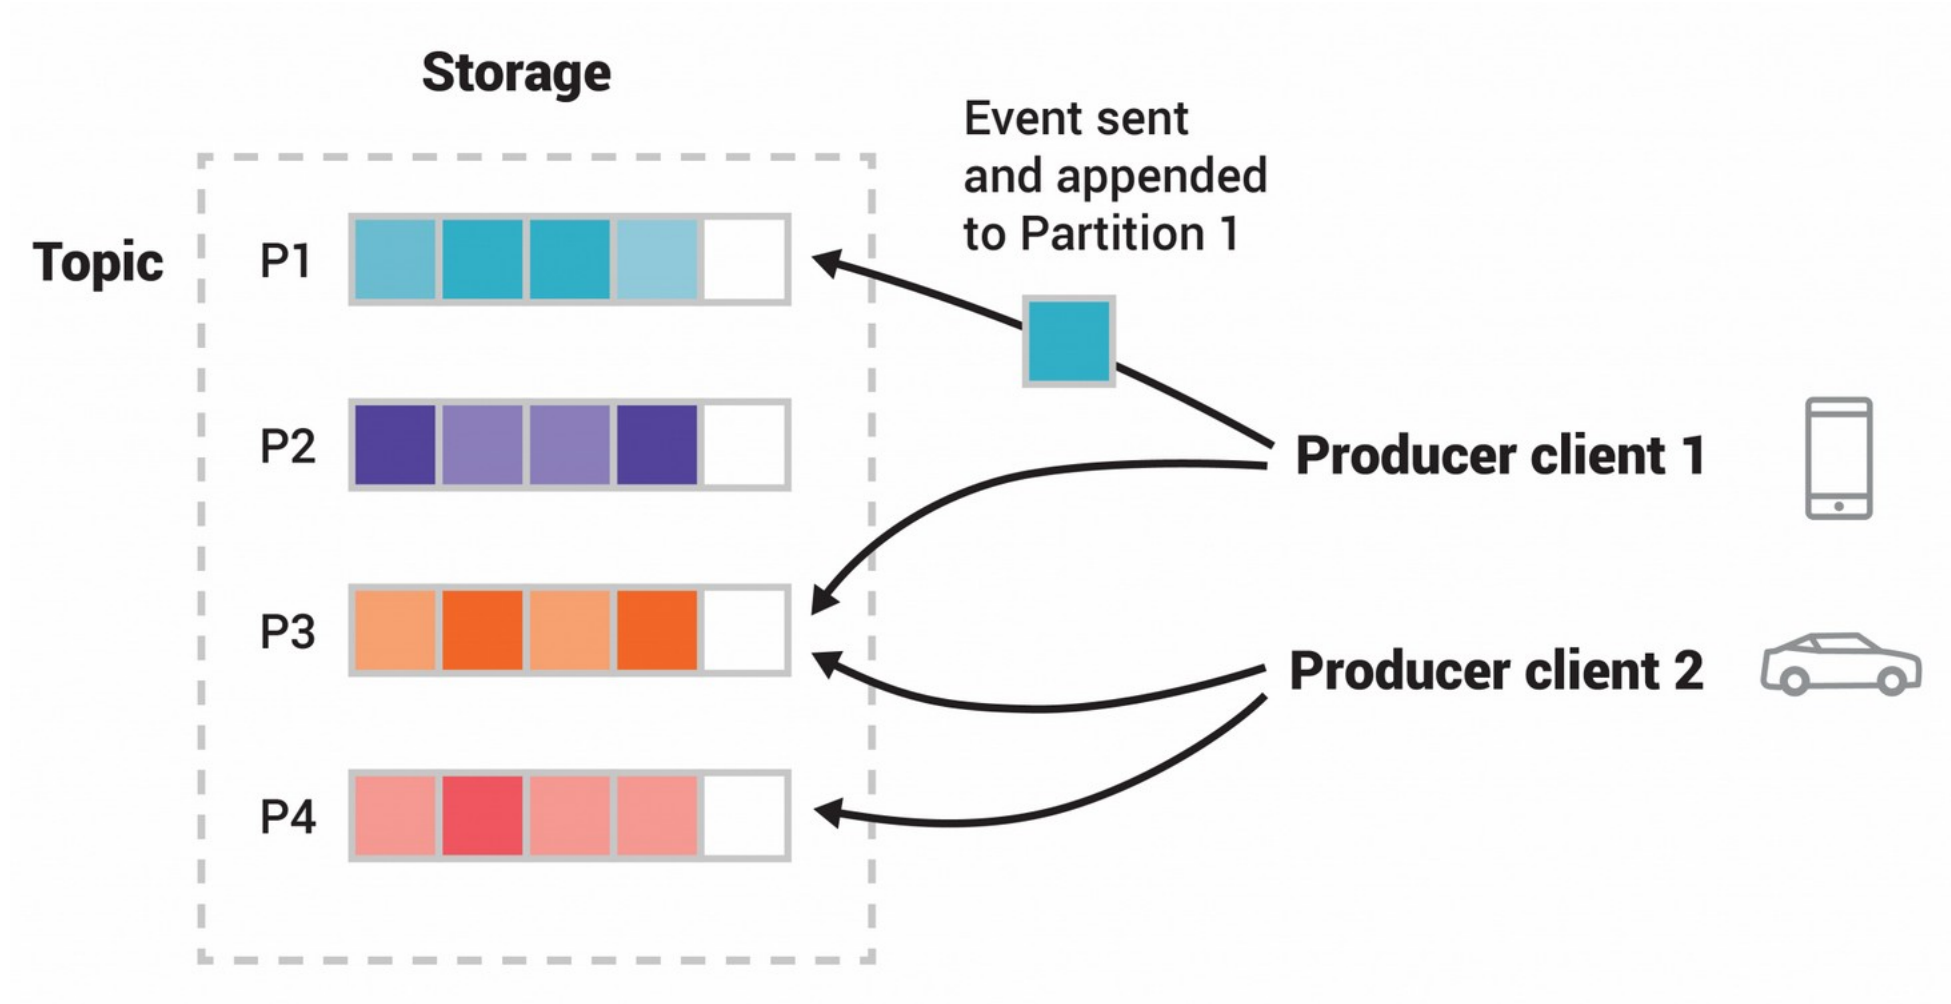
<https://kafka.apache.org/documentation>

- Event streaming is the practice of:
 - capturing data in real-time from event sources like databases, sensors, mobile devices, cloud services, and other applications in the form of streams of events;
 - storing these event streams durably for later retrieval;
 - manipulating, processing, and reacting to the event streams in real-time as well as retrospectively; and
 - routing the event streams to different destination technologies as needed.
- Ensures a continuous flow and interpretation of data.
- Crucial for data platforms, event-driven architectures, and microservices.
- Better throughput, built-in partitioning, replication, and fault-tolerance which makes it a good solution for large scale message processing applications.

Key Concepts

- **Server:** *Brokers*: form the storage layer; *Connectors*: continuously import and export data as event streams; A server *cluster* should be highly scalable and fault-tolerant.
- **Client:** *Producers* and *Consumers*; Fully *decoupled* and *agnostic* of each other
- **Event:** An event has a key, value, timestamp, and optional metadata headers; (also called *Record* or *Message*);
- **Topic:** Multi-producer and multi-subscriber; Events are not deleted after consumption by default.
- **Partition:** A topic is spread over a number of "buckets" located on different brokers → important for scalability.
- **Replication:** There are always multiple brokers that have a copy of the data in topics → fault-tolerance and high-availability

Partitions



Events with the same event key (e.g., a customer or vehicle ID) are written to the same partition, and Kafka guarantees that any consumer of a given topic-partition will always read that partition's events in exactly the same order as they were written.

Use Cases

Real-time, high volume, data centralization

- Traditional Messaging
- Website Activity Tracking
- Operational monitoring data (Metrics)
- Log Aggregation
- Stream Processing
 - Multiple stage processing pipelines
- Event Sourcing
 - state changes are logged as a time-ordered sequence of records
- Commit Log, etc.

Kafka API Example (Producer)

```
static final String SERVER = "localhost:9092";
static final String TOPIC = "myTopic";

Properties props = new Properties();
props.setProperty(ProducerConfig.BOOTSTRAP_SERVERS_CONFIG,
                  SERVER);
props.put(ProducerConfig.KEY_SERIALIZER_CLASS_CONFIG,
          StringSerializer.class.getName());
props.put(ProducerConfig.VALUE_SERIALIZER_CLASS_CONFIG,
          StringSerializer.class.getName());

try (Producer<String, String> producer =
    new KafkaProducer<>(props)) {
    ProducerRecord record =
        new ProducerRecord(TOPIC, "Hey Kafka!");
    producer.send(record);
    producer.flush(); producer.close();
}
```

Serialization is the process of converting an object into a stream of bytes that are used for transmission.

Kafka API Example (Consumer)

```
Properties props = new Properties();
props.setProperty(ConsumerConfig.BOOTSTRAP_SERVERS_CONFIG, SERVER);
props.put(ConsumerConfig.GROUP_ID_CONFIG, "KafkaExampleConsumer");
props.put(ConsumerConfig.KEY_DESERIALIZER_CLASS_CONFIG,
    StringDeserializer.class.getName());
props.put(ConsumerConfig.VALUE_DESERIALIZER_CLASS_CONFIG,
    StringDeserializer.class.getName());

try (Consumer consumer = new KafkaConsumer(props)) {

    consumer.subscribe(Collections.singletonList(TOPIC));

    ConsumerRecords<Long, String> consumerRecords
        = consumer.poll(Duration.ofSeconds(1));

    for (ConsumerRecord record : consumerRecords) {
        System.out.printf("Consumer Record: (%d, %s, %d, %d)\n",
            record.key(), record.value(),
            record.partition(), record.offset());
    };

    ///message should not be considered as consumed
    consumer.commitAsync();
```


MicroProfile Reactive Messaging

<https://github.com/eclipse/microprofile-reactive-messaging>

- Provides an easy and **standard** way to asynchronously send, receive, and process messages that are received as *continuous streams of events*.
- Provides a *Connector API* so that your methods can be connected to external messaging systems that produce and consume the streams of events, such as Apache Kafka.
- Implementations: Lightbend Alpakka, SmallRye Reactive Messaging, Open Liberty, Quarkus

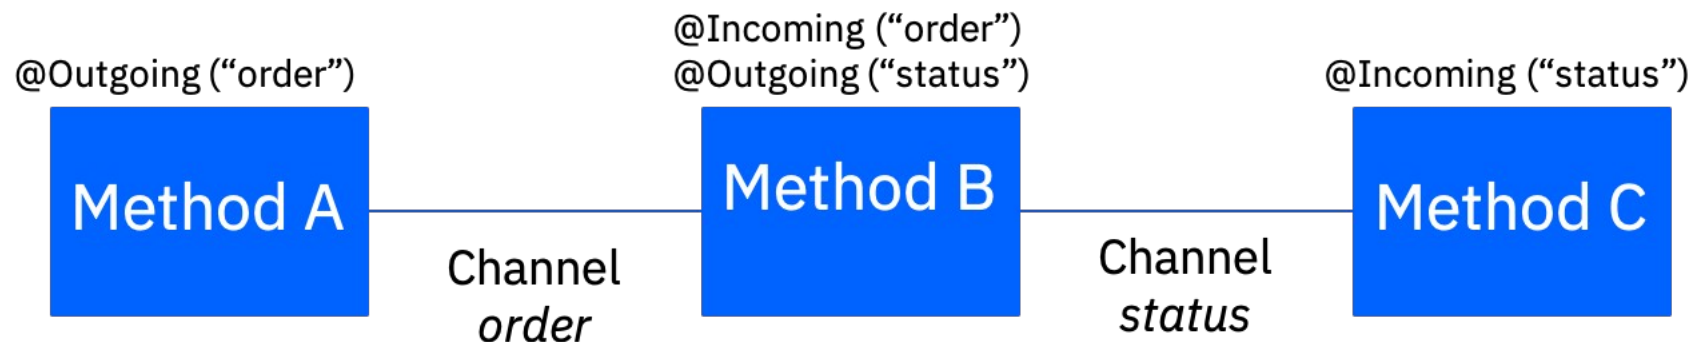
Architecture

- An application using Reactive Messaging is composed of CDI beans **consuming**, **producing** and **processing** messages. These messages can be internal to the application or can be sent and received via different message brokers.
- A **channel** is a name indicating which source or destination of messages is used.
 - Internal channels are local to the application. They allow implementing multi-step processing where several beans from the same application form *a chain of processing*.
 - Channels can be connected to remote brokers or various message transport layers such as Apache Kafka or to an AMQP broker. These channels are managed by connectors.
- A **message** is an envelope wrapping a payload. A message is sent to a specific channel and, when received and processed successfully, acknowledged.

Reactive Messaging API

<https://download.eclipse.org/microprofile/microprofile-reactive-messaging-1.0/apidocs/>

- *@Incoming*: Used to signify a subscriber to incoming messages.
- *@Outgoing*: Used to signify a publisher of outgoing messages.



<https://download.eclipse.org/microprofile/microprofile-reactive-messaging-1.0/microprofile-reactive-messaging-spec.pdf>

Example (Producer)

```
@ApplicationScoped
public class ProducerBean{
    private Random random= new Random();

    @Outgoing("bar")
    public Flowable<Integer> generatePrices() {
        return Flowable.interval(5, TimeUnit.SECONDS)
            .map(tick -> {
                int price = random.nextInt(1000);
                return price;
            });
    }
}
```

The `io.reactivex.rxjava3.core.Flowable` class implements the Reactive Streams Publisher Pattern and offers factory methods, intermediate operators and the ability to consume reactive dataflows.

Example (Consumer)

```
@ApplicationScoped
public class ConsumerBean {

    @Incoming ("foo")
    public void consume(int price) {
        System.out.println(
            "Consumer recieved: " + price
            + " @" + System.currentTimeMillis());
    }

}
```

microprofile-config.properties

```
mp.messaging.outgoing.bar.connector=liberty-kafka
mp.messaging.outgoing.bar.bootstrap.servers=localhost:9092
mp.messaging.outgoing.bar.topic = myTopic
mp.messaging.outgoing.bar.key.serializer=org.apache.kafka.common.serialization.StringSerializer
mp.messaging.outgoing.bar.value.serializer=org.apache.kafka.common.serialization.IntegerSerializer
```

```
mp.messaging.incoming.foo.connector=liberty-kafka
mp.messaging.incoming.foo.bootstrap.servers=localhost:9092
mp.messaging.incoming.foo.topic = myTopic
mp.messaging.incoming.foo.group.id=foo-reader
mp.messaging.incoming.foo.key.deserializer=org.apache.kafka.common.serialization.StringDeserializer
mp.messaging.incoming.foo.value.deserializer=org.apache.kafka.common.serialization.IntegerDeserializer
```

Dependencies

```
<feature>microProfile-3.0</feature>  
<feature>mpReactiveMessaging-1.0</feature>
```

```
<dependency>  
  <groupId>org.eclipse.microprofile.reactive.messaging</groupId>  
  <artifactId>microprofile-reactive-messaging-api</artifactId>  
  <version>1.0</version>  
  <scope>provided</scope>  
</dependency>
```

```
<dependency>  
  <groupId>org.apache.kafka</groupId>  
  <artifactId>kafka-clients</artifactId>  
  <version>2.8.1</version>  
</dependency>
```

```
<dependency>  
  <groupId>io.reactivex.rxjava3</groupId>  
  <artifactId>rxjava</artifactId>  
  <version>3.1.2</version>  
</dependency>
```